

assignment-5-bhs-atharva-pingale

April 9, 2024

1. Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset

Reference <https://inside-machinelearning.com/en/how-to-do-linear-regression-with-keras/>

```
[ ]: import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import pandas as pd
```

```
[ ]: data=pd.read_csv("/content/housing (1).csv")
data
```

```
[ ]:
      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD  TAX  \
0    0.00632  18.0    2.31    0  0.538  6.575  65.2  4.0900    1  296
1    0.02731   0.0    7.07    0  0.469  6.421  78.9  4.9671    2  242
2    0.02729   0.0    7.07    0  0.469  7.185  61.1  4.9671    2  242
3    0.03237   0.0    2.18    0  0.458  6.998  45.8  6.0622    3  222
4    0.06905   0.0    2.18    0  0.458  7.147  54.2  6.0622    3  222
..      ...    ...    ...    ...    ...    ...    ...    ...
501  0.06263   0.0   11.93    0  0.573  6.593  69.1  2.4786    1  273
502  0.04527   0.0   11.93    0  0.573  6.120  76.7  2.2875    1  273
503  0.06076   0.0   11.93    0  0.573  6.976  91.0  2.1675    1  273
504  0.10959   0.0   11.93    0  0.573  6.794  89.3  2.3889    1  273
505  0.04741   0.0   11.93    0  0.573  6.030  80.8  2.5050    1  273

      PTRATIO      B  LSTAT  MEDV
0         15.3  396.90   4.98  24.0
1         17.8  396.90   9.14  21.6
2         17.8  392.83   4.03  34.7
3         18.7  394.63   2.94  33.4
4         18.7  396.90   5.33  36.2
..      ...    ...    ...    ...
501        21.0  391.99   9.67  22.4
502        21.0  396.90   9.08  20.6
```

```
503      21.0  396.90   5.64  23.9
504      21.0  393.45   6.48  22.0
505      21.0  396.90   7.88  11.9
```

```
[506 rows x 14 columns]
```

```
[ ]: data.dtypes
```

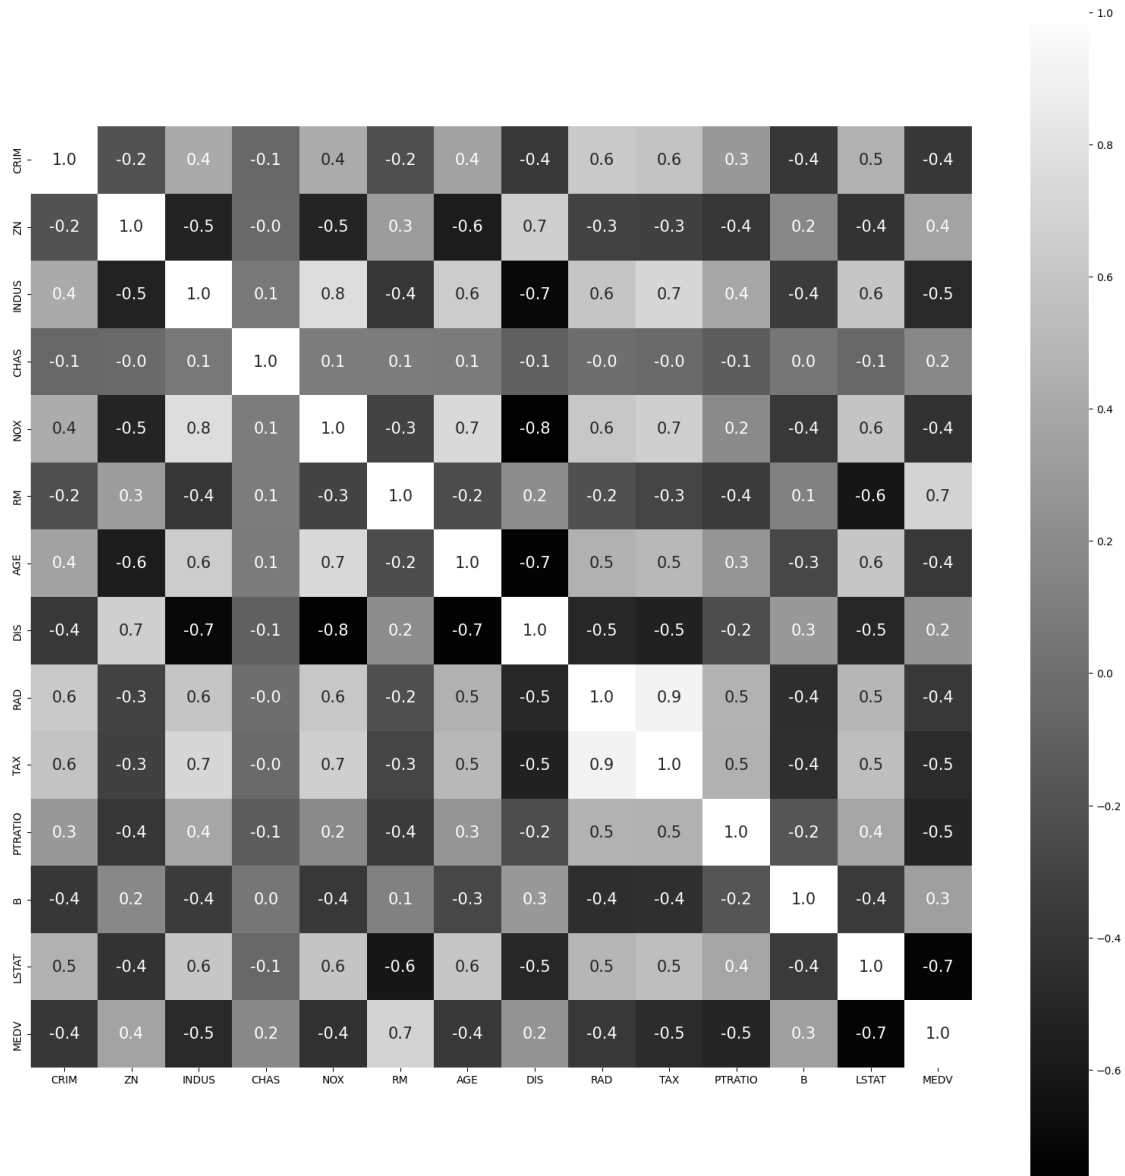
```
[ ]: CRIM      float64
      ZN       float64
      INDUS   float64
      CHAS     int64
      NOX      float64
      RM       float64
      AGE      float64
      DIS      float64
      RAD      int64
      TAX      int64
      PTRATIO  float64
      B        float64
      LSTAT    float64
      MEDV     float64
dtype: object
```

```
[ ]: # Finding out the correlation between the features
      corr = data.corr()
      corr.shape
```

```
[ ]: (14, 14)
```

```
[ ]: # Plotting the heatmap of correlation between features
      plt.figure(figsize=(20,20))
      sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True,
                  ↪annot_kws={'size':15}, cmap='gray')
```

```
[ ]: <Axes: >
```



```
[ ]: X =data.drop(['MEDV'], axis = 1)# data['area']#
y = data['MEDV']
```

```
[ ]: X
```

```
[ ]:
      CRIM    ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD  TAX  \
0   0.00632  18.0   2.31    0  0.538   6.575   65.2   4.0900    1  296
1   0.02731   0.0   7.07    0  0.469   6.421   78.9   4.9671    2  242
2   0.02729   0.0   7.07    0  0.469   7.185   61.1   4.9671    2  242
3   0.03237   0.0   2.18    0  0.458   6.998   45.8   6.0622    3  222
4   0.06905   0.0   2.18    0  0.458   7.147   54.2   6.0622    3  222
```

```

..      ...      ...      ...      ...      ...      ...      ...      ...
501  0.06263      0.0  11.93      0  0.573  6.593  69.1  2.4786      1  273
502  0.04527      0.0  11.93      0  0.573  6.120  76.7  2.2875      1  273
503  0.06076      0.0  11.93      0  0.573  6.976  91.0  2.1675      1  273
504  0.10959      0.0  11.93      0  0.573  6.794  89.3  2.3889      1  273
505  0.04741      0.0  11.93      0  0.573  6.030  80.8  2.5050      1  273

```

```

      PTRATIO      B  LSTAT
0      15.3  396.90   4.98
1      17.8  396.90   9.14
2      17.8  392.83   4.03
3      18.7  394.63   2.94
4      18.7  396.90   5.33
..      ...      ...      ...
501     21.0  391.99   9.67
502     21.0  396.90   9.08
503     21.0  396.90   5.64
504     21.0  393.45   6.48
505     21.0  396.90   7.88

```

[506 rows x 13 columns]

```
[ ]: X.columns
```

```
[ ]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
          'PTRATIO', 'B', 'LSTAT'],
          dtype='object')
```

```
[ ]: y
```

```
[ ]: 0      24.0
     1      21.6
     2      34.7
     3      33.4
     4      36.2
     ...
    501     22.4
    502     20.6
    503     23.9
    504     22.0
    505     11.9
    Name: MEDV, Length: 506, dtype: float64
```

```
[ ]: y.describe()
```

```
[ ]: count      506.000000
     mean       22.532806
```

```
std          9.197104
min           5.000000
25%          17.025000
50%          21.200000
75%          25.000000
max          50.000000
Name: MEDV, dtype: float64
```

```
[ ]: # Splitting to training and testing data

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3,
↳random_state = 4)
```

```
[ ]: from sklearn.preprocessing import MinMaxScaler
# Instantiate the scaler and fit to training dataset, X_train
scaler = MinMaxScaler()
scaler.fit(X_train)

# Replace unscaled values with scaled values
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
[ ]:
```

```
[ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
```

```
[ ]: model = Sequential()L_Assignment2
model.add(Dense(64, input_dim =13, activation = 'relu'))
model.add(Dropout(0.15))
model.add(Dense(64, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(40, activation = 'relu'))
model.add(Dropout(0.15))
model.add(Dense(54, activation = 'relu'))
model.add(Dropout(0.18))
model.add(Dense(1))
```

```
[ ]: model.compile(optimizer = 'adam', loss = 'mean_squared_error', metrics =
↳['mse', 'mae'])
```


```
[ ]: history = model.fit(X_train, y_train, validation_split=0.2, epochs=200)
```

```
[ ]: import matplotlib.pyplot as plt
```

```
#plot the loss and validation loss of the dataset
plt.plot(history.history['mae'], label='mae')
plt.plot(history.history['val_mae'], label='val_mae')

plt.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7f91acbf7f70>
```



```
[ ]:
```

```
[ ]: scores = model.evaluate(X_test, y_test, verbose = 0)

print('Mean Squared Error : ', scores[1])
print('Mean Absolute Error : ', scores[2])
```

```
Mean Squared Error : 17.62657928466797
Mean Absolute Error : 2.6688265800476074
```

```
[ ]: Y_pred = model.predict(X_test)
Y_pred
```


```
[ ]: from sklearn.metrics import r2_score  
  
print('r2 score: ', r2_score(y_test,Y_pred))
```

r2 score: 0.831195498482862

```
[ ]: Y_pred = model.predict(X_test)  
  
a = plt.axes(aspect='equal')  
  
plt.xlabel('True values')  
plt.ylabel('Predicted values')  
plt.xlim([0, 50])  
plt.ylim([0, 50])  
plt.plot([0, 50], [0,50])  
plt.scatter(y_test,Y_pred)  
plt.plot()
```

5/5 [=====] - 0s 3ms/step


```
[ ]: []
```



```
[ ]:
```

```
[ ]: plt.plot(y_test)
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f91b0719c70>]
```



```
[ ]: plt.plot(Y_pred)
      Y_pred
```

```
[ ]: print(Y_pred[:5])
      print(y_test[:5])
      y_test.head()
```

```
[[16.747326]
 [24.475298]
 [19.15377 ]
 [16.203077]
 [42.90115 ]]
8      16.5
289    24.8
68     17.4
211    19.3
```



```
226      37.6
Name: MEDV, dtype: float64
```

```
[ ]: 8      16.5
      289    24.8
      68     17.4
      211   19.3
      226    37.6
Name: MEDV, dtype: float64
```

Using ML MODEL lm for Linear Regression

```
[ ]: # Import library for Linear Regression
     from sklearn.linear_model import LinearRegression
```

```
[ ]: # Create a Linear regressor
     lm = LinearRegression()

     # Train the model using the training sets
     lm.fit(X_train, y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: # Value of y intercept
     lm.intercept_
```

```
[ ]: 27.380936280008473
```

```
[ ]: #Converting the coefficient values to a dataframe
     coefficients = pd.DataFrame([X.columns,lm.coef_]).T
     coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coefficients'})
     coefficients
```

```
[ ]:  Attribute Coefficients
0      CRIM   -10.90502
1       ZN     5.56777
2     INDUS  -0.238526
3      CHAS    4.693448
4      NOX   -7.01579
5       RM   15.865749
6      AGE   -0.33478
7      DIS  -17.068774
8      RAD    7.503741
9      TAX   -7.370872
10   PTRATIO  -7.550784
11       B     3.709485
12   LSTAT  -18.452583
```

```
[ ]: # Model prediction on train data
y_pred = lm.predict(X_train)
```


```
[ ]: # Model Evaluation
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.089861094971128
MSE: 19.073688703469028
RMSE: 4.367343437774161
```

```
[ ]: # Visualizing the differences between actual prices and predicted values
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



```
[ ]: # Checking residuals
plt.scatter(y_pred,y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



```
[ ]: # Checking Normality of errors
sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```


<ipython-input-57-c62ec83682b5>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(y_train-y_pred)
```



```
[ ]: # Predicting Test data with the model
y_test_pred = lm.predict(X_test)
```

```
[ ]: # Model Evaluation
acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred)) * (len(y_test)-1) / (len(y_test)-X_test.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7121818377409196
Adjusted R^2: 0.6850685326005714
MAE: 3.8590055923707407
MSE: 30.05399330712412
RMSE: 5.482152251362973
```